House of the Rising Sun

Strategically Loyal Republicans and Prisoner's-Dilemma Democrats

Da Gong* Andong Yan† Cantay Caliskan‡ Ezgi Kibris§ Eunseong Oh¶

May 2026

Pilot draft prepared for EIP 2026

Abstract

A serious concern about cult of personality around a party leader is that the legislature stops checking the executive and instead defers to the leader — whether because members fear punishment from party leadership for deviating from the party line, or because constituents themselves come to treat the leader's signal as their dominant preference. We study this dynamic in the contemporary United States, where Donald Trump is an unprecedentedly salient figure whose social media posts rapidly structure elite attention and partisan messaging. When a salient leader initiates a novel stance or shifts position, how do members of Congress respond, and to what extent does the response depend on the leader's cue versus constituent preferences? Treating Trump's Truth Social posts as datable coordination signals, we measure congressional reactions in a near-universe of sitting members' tweets (2022–2026), classifying *relevance* and *stance* at scale with open-weight large language models (LLMs). Across 24 high-salience signals, we document a large and stable partisan stance gap alongside wide variation in within-Republican consensus — from near-unanimity on tariffs (84% support) to an even split on Ukraine (49%). Using regression discontinuity (RD) in event-time around three sharp signals, we find *issue-specific strategic loyalty/silence*: on the controversial Ukraine signals, only Democrats raise topic attention (+3 to +6 percentage points), while Republicans as a group do not move — yet the Republicans who do speak shift toward Trump's framing (blame 0.21; reversal +0.46, conditional on engagement). On the consolidated tariff signal, both parties raise attention and Republicans also shift expressed stance (+0.18). The pattern is consistent with a loyalty–accountability trade-off in which cross-pressured in-party members manage risk by going quiet rather than by opposing. For the out-party, individual incentives to attack the leader undercut the collectively optimal strategy of denying him salience — a *prisoner's dilemma* in counter-signaling. We view this study as establishing the measurement pipeline and the core substantive pattern: expansion to the full event space, classifier validation, and cross-validation against roll-call voting are ongoing.

*SUNY Geneseo. Corresponding author: `dgong005@ucr.edu`.

†University of Hong Kong.

‡University of Rochester.

§Rochester Institute of Technology.

¶University of California, Riverside.

1 Introduction

A central concern in the study of democratic backsliding is that a cult of personality around a party leader erodes the separation of powers: the legislature stops checking the executive and instead defers to the leader (Levitsky and Ziblatt, 2018; Graham and Svolik, 2020). One channel operates directly through representation. A member of Congress may fail to reflect local constituents' preferences for two distinct reasons: either the member fears punishment from party leadership for deviating from the party line, or constituents themselves come to treat the leader's signal as dominant over their issue-specific preferences, such that adhering to the leader is, in their view, a form of representation. Donald Trump is an unprecedentedly salient figure in American politics, and his social media account issues high-salience signals that rapidly structure elite attention and partisan messaging. That mechanism is also an empirical opportunity: a stream of datable, leader-issued signals lets us ask *When a salient party leader shifts position, how do members of Congress respond, and to what extent does that response depend on the leader's cue versus constituent preferences?*

The question sits within the study of *representation*—whether a legislator's public behavior tracks the preferences of those they represents. Classic accounts locate representation in the correspondence between constituents' positions and a member's record (Miller and Stokes, 1963; Clinton, 2006), and recent work broadens it from specific policy positions to the deeper values and worldviews constituents hold, measuring district opinion at scale with large surveys (Vishwanath, 2026; Ansolabehere and Kuriwaki, 2022, 2025). Our setting introduces a competing principal into this picture: a salient national leader whose signal can override the local district as the object a member answers to. Two further literatures sharpen the resulting tension. The first treats legislators as strategic users of presidential and elite cues who tend to *follow* the issue attention of their supporters more than they lead it (Noble, 2023; Barberá et al., 2019), organized ultimately around reelection (Fenno, 1973; Mayhew, 1974). The second is electoral accountability: misalignment with district preferences reduces vote share (Canes-Wrone et al., 2002), and partisan loyalty on divisive issues can be electorally costly (Carson et al., 2010; Hirano et al., 2010). Representing the *local* district and deferring to the *national* leader can pull apart: following the leader is attractive but not free—it trades partisan coordination and base-signaling benefits against the electoral cost of getting out of step with the district. How a member resolves that trade-off should depend on whether the leader's position is *locally controversial*.

We argue—and find—that the most consequential margin of adjustment is often not *what* a member says but *whether they speak at all*. A cross-pressured Republican facing a controversial Trump signal can align publicly (and alienate constituents), oppose openly (and invite primary backlash), or **say nothing**. Strategic silence is frequently optimal, and treating it as missing data systematically understates local pressure, because the members under the most pressure are precisely the ones who disappear from the text. We therefore keep two behavioral margins separate throughout: *engagement / attention*—whether and how much a member speaks on the topic—and *stance*—the position taken conditional on speaking.

This paper. We assemble a near-universe of original tweets by sitting members of Congress (2022–January 2026), identify the precise Truth Social post that constitutes each Trump signal, and classify member tweets for relevance and stance using open-weight LLMs (Qwen-2.5) in a reproducible pipeline. We then (i) describe the partisan stance landscape across 24 high-salience signals, which reveals a spectrum from consolidated to controversial Republican positions, and (ii) estimate the causal effect of three sharp signals using *regression discontinuity in event-time* (Hausman and Rapson, 2018) around the post timestamp T_e.

Findings. Three results stand out. First, the partisan stance gap is large and remarkably stable: Democrats oppose Trump's position in most relevant tweets on nearly every issue, while Republican support ranges from 84% (tariffs) down to 49% (Ukraine) and lower, tracing a clear consolidated-to-controversial spectrum. Second, on the *controversial* Ukraine signals, the within-member jump in topic attention at T_e is driven entirely by Democrats; Republicans as a group do not move—the signature of strategic silence—while the Republicans who *do* engage shift their stance toward Trump's framing. Third, on the *consolidated* tariff signal both parties raise attention and Republicans additionally move their expressed stance toward Trump. The contrast—silence on the controversial issue, full mobilization on the consolidated one—is the paper's central object.

Contribution. We offer (i) a scalable, reproducible LLM pipeline for measuring elite rhetorical attention and stance across millions of candidate tweets; (ii) a timing-based identification of how a dominant national signaler moves congressional behavior; and (iii) evidence that strategic silence is *issue-specific*, reframing congressional "silence" from missing data into a measurable strategic act. This is a pilot study; classifier validation and the full identification battery are in progress (Section 7).

2 Theoretical framework and predictions

The trade-off. Legislators use presidential cues strategically (Noble, 2023), tend to follow rather than lead supporter attention (Barberá et al., 2019), and invest effort where it is electorally productive (Fouirnaies and Hall, 2022), with the in-party rallying around a co-partisan president while the out-party resists (Lebo and O'Geen, 2011; Cohen, 1995; Eshbaugh-Soha and Peake, 2011). Against this pull, misalignment with the district is punished (Canes-Wrone et al., 2002), especially partisan loyalty on divisive votes (Carson et al., 2010). The resolution should depend on whether the leader's position is locally controversial.

What counts as representation failure—and our assumptions. Our normative benchmark is the *trustee*, or gatekeeping, conception of representation: a member should exercise independent judgment in the service of constituents' *interests* rather than merely transmit their *will*—least of all a low-information affective cue (the Burkean interest-versus-will distinction; Miller and Stokes, 1963; Canes-Wrone et al., 2001; Arnold, 1990). By this benchmark, abandoning a district's issue preference to follow the presidential cue is a representation *failure*, and it is a failure *whichever* channel drives the following—whether the district itself has come to demand the cue (a voter-preference channel) or the member fears the leader's punishment (a loyalty/career channel). Both substitute the leader's signal for the representative's own job.

We make this claim under deliberately weak assumptions. We take voters to be rational *at the policy level*—able to evaluate a concrete policy—without assuming their policy preferences are good for them; overweighting the presidential cue can lead voters to act *against* the policy interests they would otherwise hold, harming themselves. We do *not* assume that representatives can observe constituents' true preferences, which is often infeasible (reviewed in Vishwanath, 2026); we assume only that observed behavior is an *equilibrium* that members believe maximizes their reelection prospects. A natural objection is that voters are sometimes not even policy-level rational—e.g., a tariff-exposed agricultural district convinced new tariffs will help it when they will not. That only *deepens* the failure (the member now follows a mistaken will against a harmed interest), and so strengthens rather than weakens the argument.

The implication for our design is that presidential-cue-following is a representation failure *regardless of channel*—which is what the empirical strategy can establish. The open challenge is to *disentangle the channels*—voter demand versus loyalty pressure—which we sketch in the discussion and complete after this draft.

Controversial vs. consolidated. We call an issue *controversial* when opinion is internally divided—within the party's own electorate and between the party line and the district—and *consolidated* when the party and the district agree. Empirically, the within-party distribution of expressed stance is a useful proxy: tariffs, where Republicans are near-unanimous, sit at the consolidated pole; Ukraine, where Republicans split roughly evenly, sits at the controversial pole.[1]

Two asymmetric traps. The trade-off plays out differently for the two parties, and the difference is the paper's organizing theme. For **Republicans**, the leader is a co-partisan whose displeasure is electorally dangerous, so the binding problem is *strategic loyalty*: on a controversial signal a cross-pressured member cannot safely align (alienating the district) or oppose (inviting a primary), so the optimal move is to go quiet—loyalty enforced through silence. For **Democrats** the signal comes from the out-party, and the trap is a collective one resembling a *prisoner's dilemma*. Each Democrat has a private incentive to attack a controversial Trump post—it rallies the base and contests his framing—but loud, coordinated opposition *amplifies* the issue's salience on Trump's preferred terms, which is collectively costly to the party even as it is individually rational. Individually optimal defection (engage) yields a collectively worse outcome (amplify); the cooperative outcome (cede the floor, deny salience) is unstable. The two traps generate opposite engagement signatures on controversial issues: Republican *under*-engagement (silence) and Democratic *over*-engagement

[1] In the full design this classification is set *ex ante* from Cooperative Election Study (CES) district preferences and a divisiveness measure (Ansolabehere and Kuriwaki, 2022, 2025); here we use the observed partisan stance distribution as a transparent, if endogenous, proxy and flag the CES-based version as the confirmatory step.

(amplification).

Predictions. The framework implies an asymmetry by party and issue type:

- **H1 (controversial, extensive margin).** On a controversial signal, out-party (Democratic) engagement jumps at the post while in-party (Republican) engagement does not—the signature of strategic loyalty/silence.
- **H2 (consolidated, extensive margin).** On a consolidated party-agenda signal, *both* parties' engagement jumps.
- **H3 (intensive margin under selection).** Conditional on engaging, Republicans who speak move toward Trump's stated position, concentrated among members whose districts impose the least local cross-pressure.
- **H4 (the Democratic dilemma).** Democratic engagement with a controversial signal does not retreat in the districts where amplification is most costly; if anything it is strongest there (e.g., low-Trump, high-local-salience districts)—the signature of individually rational opposition that is collectively self-amplifying, rather than party-disciplined restraint.

Republican silence is the public-messaging analogue of the classic logic of position avoidance and ambiguity under cross-pressure (Shepsle, 1972; Arnold, 1990; Mayhew, 1974); the Democratic dilemma is the agenda-amplification trap a dominant opposing signaler creates (Barberá et al., 2019; Faris et al., 2017).

3 Data

Congressional communication. The near-universe of original tweets posted by sitting members of Congress from 2022 through January 2026, excluding retweets, replies, and quote-tweets, restricted to English, and dropping trivially short posts. Tweets are linked to members via a Bioguide crosswalk. Twitter is a defensible window onto congressional position-taking: members use it to advertise positions and engage in partisan messaging (Hemphill et al., 2013), and the issues they discuss correlate highly with their broader expressed agenda (Barberá et al., 2019). Communication is also not redundant with the roll-call record: the partisanship of speech and the polarization of votes follow different time paths (Gentzkow et al., 2019).

Trump signal corpus. Original Truth Social posts (Feb 2022–Jan 2026), HTML stripped. For each event, the specific triggering post and its exact timestamp T_e are identified by keyword search over the archive followed by researcher review, and recorded with the post ID, URL, and text—so the treatment timing is explicit and reproducible for every event.

District covariates. District-level 2020 Trump vote share (a proxy for district favor toward the national Trump agenda), CES policy-preference estimates at the district level (Ansolabehere and Kuriwaki, 2022), and standard demographics. These support the heterogeneity cuts by district Trump share and local issue sentiment.

4 Measurement: an LLM pipeline for attention and stance

For each event we define a window around T_e and query all original member tweets in it (typically 90,000–120,000). A seed dictionary—expanded by the King et al. (2017) query-expansion procedure and an LLM step—narrows the window to topical candidates using anchor-aware contextual matching. Two classification stages follow, both run locally on open-weight models with checkpoint/resume:

- **Relevance** (Qwen-2.5-7B-Instruct, 4-bit). Each candidate is labeled 0 (not relevant), 1 (indirectly relevant), or 2 (clearly about the event), with confidence from constrained single-token generation.
- **Stance** (Qwen-2.5-14B-Instruct, 4-bit). Each classified tweet receives, via constrained generation, a stance toward Trump's stated position—oppose (-1), neutral (0), or support (+1)—with "unclear" set aside.

The same protocol underlies recent text-as-data practice (Benoit et al., 2025; Laurer et al., 2024; Kawintiranon and Singh, 2022). Every classified tweet carries a stance score, which we aggregate to member-day means for the panel. Classifier validation against an expert gold standard is in progress and is discussed in Section 7.

5 Empirical strategy

Two sources of variation support two complementary designs. Most issues are not punctuated by a single dramatic post; they simmer continuously, which calls for a time-series design (Strategy 1, panel VAR). Some signals are sharp, datable switches, which call for an event-time discontinuity (Strategy 2, RD). The two are deliberately separated: continuous co-movement of attention, stance, and framing in Strategy 1; discrete causal shocks in Strategy 2. This paper's estimated results come from Strategy 2; we describe Strategy 1 because it is the companion design that recovers the *persistent* alignment our theory also predicts, and its estimates are forthcoming.

5.1 Strategy 1: persistent topics and a panel VAR

For continuously contested issues we ask a *who-leads-whom* question: when Trump shifts attention to a topic—or shifts his position on it—does Congress follow, how fast, and for whom? We treat Trump, each legislator i, and a set of $\sim$20 media figures (journalists and broadcasters) as parallel daily time series, extending the design used to measure issue-attention leadership in Congress (Barberá et al., 2019) to carry several outcome series per actor. For each actor $\times$ topic j $\times$ day t we model the log-odds of the actor's attention share $Z^A_{ijt} = \log(A_{ijt}/(1 - A_{ijt}))$ and the mean expressed stance $s_{ijt} \in [-1, 1]$ (Trump's series, s^T_{jt}, is scored from Truth Social), and fit a bivariate panel VAR with $p = 7$ daily lags and topic fixed effects α_j:

$$\begin{pmatrix} Z^A_{ijt} \\ s_{ijt} \end{pmatrix} = \alpha_j + \sum_{k=1}^{7} \mathbf{B}_{ik} \begin{pmatrix} Z^A_{ij,t-k} \\ s_{ij,t-k} \end{pmatrix} + \boldsymbol{\varepsilon}_{ijt}. \tag{1}$$

We summarize the system with cumulative impulse-response functions answering two questions: when Trump's attention to a topic jumps, how much and how fast does legislator i's attention follow; and when Trump's *stance* shifts, does s_{ijt} move toward s^T_{jt}? Carrying media actors is essential rather than cosmetic: a common media shock could drive both Trump and Congress, so the VAR treats media as a confounder to be conditioned on, not assumed away (Barberá et al., 2019). Tracking stance as well as attention is what lets Strategy 1 speak to *alignment*, not merely agenda-setting—the difference between Congress noticing what Trump raises and Congress moving toward what he says. We split the impulse responses by party and by district Trump-share and CES local-preference tiers to test whether the persistent follow pattern is concentrated where the theory expects, and—because following the leader can itself be electorally costly (Carson et al., 2010)—*who* follows is as informative as whether anyone does. Speech and votes are distinct objects with different dynamics (Gentzkow et al., 2019), so the persistent stance series is not redundant with the roll-call record. *Strategy 1 estimates are in progress and not reported here.*

5.2 Strategy 2: novel posts as shocks (RD in event-time)

We exploit *novel, unpredictable* Trump posts—a datable moment when he raises a new issue or reverses a stated position. Because a member cannot anticipate the exact moment or direction of such a switch, the post acts as a sharp break in event-time. For member i on day t, with running variable $\tilde{d}_t = t - T_e$ and outcome Y (topic attention share, or—conditional on engagement—mean stance):

$$Y_{it} \;=\; \alpha_i \;+\; \beta\,\mathbf{1}[t \geq T_e] \;+\; \gamma\,\tilde{d}_t \;+\; \delta\left(\mathbf{1}[t \geq T_e]\cdot \tilde{d}_t\right) \;+\; \varepsilon_{it}, \tag{2}$$

where α_i is a member fixed effect, the interaction lets the local-linear time trend differ on each side of the cutoff, and β—the within-member jump at the post—is the causal estimand (Hausman and Rapson, 2018). Standard errors are clustered by member; the primary window is ± 7 days, which limits the chance that another event produces a coincident jump. The identifying assumption is that, absent the post, the outcome would have evolved continuously through T_e. The timing of Trump's posts is plausibly quasi-random: he posts impulsively, off any fixed schedule, often late at night (Almond and Du, 2020). The estimand is robust to the post being endogenous in *origin*—what the design requires is only that nothing else jumps discontinuously at T_e, so a slow-moving upstream trigger is absorbed by the local trend.

Table 1: Partisan stance distribution, selected signals (relevant tweets, confidence ≥ 0.7). Entries are within-party percentages of classified relevant tweets; "unclear" is excluded, so rows need not sum to 100. Events ordered by Republican support (consolidated $\rightarrow$ controversial).

		Democrats		**Republicans**	
Signal	T_e	% oppose	% support	% oppose	% support
Tariffs (`trump_tariffs_01`)	2025-03-03	91.3	2.3	4.8	**83.7**
DEI executive order	2025-01-19	93.1	1.4	13.0	79.5
USAID shutdown	2025-02-03	93.3	0.5	13.0	77.7
DOGE establishment	2024-12-04	84.6	3.9	21.5	61.5
Greenland (push 1)	2024-12-22	97.8	0.0	13.6	60.2
Government shutdown	2025-03-13	80.1	0.5	14.7	54.5
Ukraine blame	2025-02-19	96.3	0.6	41.7	**48.9**
NATO threat	2025-03-08	91.1	0.0	18.9	42.6
Greenland (push 2)	2026-01-14	84.1	0.9	24.7	26.3
Roe v. Wade reversal	2022-06-24	69.5	3.2	35.8	22.6

6 Results

6.1 The partisan stance landscape

Table 1 reports the partisan stance distribution for ten signals spanning the spectrum (the full set of 24 events shows the same pattern). Two facts stand out. First, the **partisan gap is large and consistent**: Democrats oppose Trump's position in the large majority of relevant tweets on nearly every policy and foreign-affairs signal, and the parties essentially never converge. Second, **within-Republican consensus varies sharply across issues**, tracing a consolidated-to-controversial spectrum: Republican support runs from 83.7% on tariffs and 79.5% on the DEI order down to 48.9% on the Ukraine blame signal, 42.6% on the NATO threat, and lower still on Roe and the second Greenland push. This spectrum is exactly the cross-issue variation the theory says should govern whether in-party members can speak freely or must manage risk.

Engagement responses mirror this. Of 22 events with computed attention metrics, 9 show a post/pre daily-attention ratio ≥ 2.0 ("confirmed"), led by Ukraine blame (8.9×), the DEI order (6.6×), and the Charlie Kirk assassination (5.5×); high-baseline issues such as tariffs and the government shutdown show smaller ratios because their pre-period attention floor is already high, not because members ignore them.

6.2 Regression discontinuity: attention

Table 2 reports the within-member RD jump at T_e for the three signals with clean, datable shocks. Both Ukraine signals produce a significant jump in topic *attention share* (+3.0 pp at each event), and the tariff signal produces larger jumps for both parties. Figures 1 and 2 show the discontinuities visually.

The heterogeneity in Table 3 resolves *who* drives the attention jumps and is the core test of H1,

Table 2: Regression discontinuity in event-time: within-member jump $\hat{\beta}$ at T_e (member FE; SE clustered by member). Attention = topic tweet share, defined for all member-days; Stance = mean Qwen stance, conditional on ≥ 1 topical tweet. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Signal / outcome	$\hat{\beta}$	SE	p	N
Ukraine blame (Feb 19, 2025) — controversial				
Attention share	+0.030	0.014	0.037**	3,922
Stance (pooled)	−0.075	0.083	0.370	803
Ukraine reversal (Sep 23, 2025) — controversial				
Attention share	+0.030	0.011	0.007***	3,911
Stance (pooled)	+0.252	0.153	0.102	589
Tariffs (CCM announcement, Feb 1, 2025) — consolidated				
Attention, Republicans	+0.130	0.030	0.000***	2,026
Attention, Democrats	+0.074	0.020	0.000***	3,255
Stance, Republicans	+0.179	0.059	0.003***	714
Stance, Democrats	+0.067	0.078	0.387	800

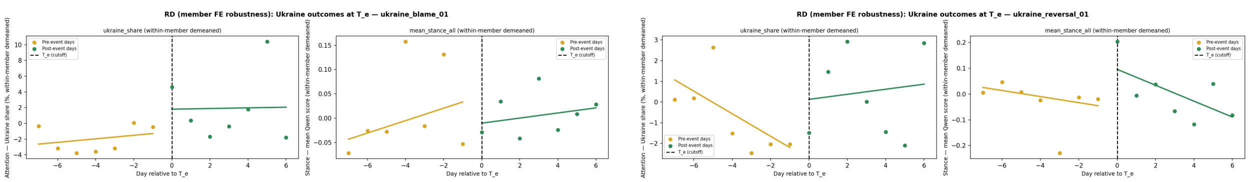

Figure 1: RD in event-time, Ukraine signals (attention share). Within-member demeaned daily means; local-linear fits pre (gold) and post (green) the post timestamp T_e. Left: blame ($\hat{\beta} = +0.030$, $p = 0.037$). Right: reversal ($\hat{\beta} = +0.030$, $p = 0.007$).

H2, and H4. On both Ukraine signals the pooled attention jump is **entirely Democratic**: Republicans show no significant response (+0.020, $p = 0.32$ on blame; +0.011, $p = 0.46$ on reversal)—the *strategic-loyalty* signature, the in-party going quiet on the controversial issue. Democrats, by contrast, jump significantly, and—consistent with the *prisoner's-dilemma* reading (H4)—the jump is *strongest* precisely where amplification is most costly to the party: low-Trump, high-local-support (pro-Ukraine) districts (+0.060, $p = 0.002$ on reversal). Democratic engagement intensifies rather than retreats under cross-pressure, the opposite of party-disciplined restraint. On the consolidated tariff signal, by contrast, *both* parties raise attention strongly and Republicans respond uniformly regardless of local tariff sentiment (+13 pp either way).

6.3 Regression discontinuity: stance

The pooled Ukraine stance RD is insignificant (Table 2), but the *conditional-on-engagement* pattern is informative and consistent with H3. Among engaged Republicans, stance moves toward Trump's framing on both events: −0.208 ($p = 0.035$) on blame—and −0.279 ($p = 0.058$) for high-Trump, low-local-support districts—and +0.463 ($p = 0.045$) on the reversal. Democrats hold their pro-Ukraine stance throughout (no significant shift). On the consolidated tariff signal, Republicans move pro-tariff (+0.179, $p = 0.003$; strongest in high-Trump, pro-tariff districts at +0.227, $p =$

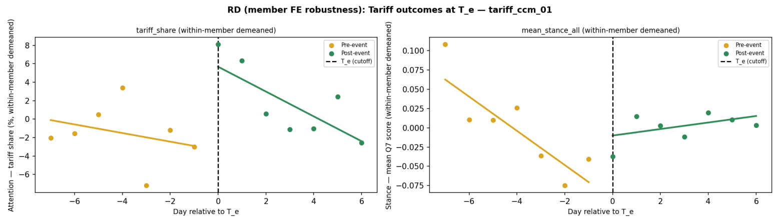

Figure 2: RD in event-time, tariff (CCM) signal (tariff-topic share). Both parties show a clear upward jump at T_e.

Table 3: RD heterogeneity, attention share, by party and local issue sentiment (member FE; SE clustered by member). "High/Low Trump" = district 2020 Trump vote tier; "Local" = district issue sentiment. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Subgroup	$\hat{\beta}$	SE	p	N
Ukraine blame — attention				
Republican (all)	+0.020	0.019	0.315	1,893
Democrat (all)	+0.042	0.021	0.051*	2,029
Democrat | Low Trump	+0.040	0.021	0.059*	1,985
Ukraine reversal — attention				
Republican (all)	+0.011	0.015	0.463	1,873
Democrat (all)	+0.047	0.016	0.004***	2,038
Democrat | Low Trump	+0.054	0.016	0.001***	1,997
Democrat | Low Trump, High Local	+0.060	0.019	0.002***	1,286
Tariffs (CCM) — attention				
Republican (all)	+0.130	0.030	0.000***	2,026
Republican | High Trump, Pro-Tariff Local	+0.118	0.042	0.006***	1,115
Republican | High Trump, Anti-Tariff Local	+0.131	0.052	0.014**	701
Democrat (all)	+0.074	0.020	0.000***	3,255
Democrat | Low Trump, Anti-Tariff Local	+0.115	0.029	0.000***	1,261

0.002), while Democrats increase engagement to oppose but do not move their mean stance.

6.4 The contrast: controversial vs. consolidated

The two cases summarize the argument. **Ukraine (controversial; 49% Republican support).** Republican districts are split, so engaging publicly carries electoral risk. The data show Democrats driving the attention jump while Republicans stay silent—and the Republicans who do speak are disproportionately the least locally constrained, and they align with Trump. **Tariffs (consolidated; 84% Republican support).** The position is core party agenda, so the local-accountability penalty for engaging is small. Both parties mobilize, and Republicans move both attention and stance toward Trump regardless of local sentiment. Strategic silence is thus *issue-specific*: the loyalty/accountability trade-off suppresses in-party engagement only when the leader's position is locally controversial.

7 Limitations and next steps

Stance measurement and validation. The stance classifier's off-the-shelf accuracy is ≈70%, and classifier error attenuates stance estimates toward zero (Burnham, 2025). A full validation is in progress: a human gold-standard with a written codebook, a stronger production classifier benchmarked per issue, and *design-based supervised learning* (Egami et al., 2024) to carry the labeling error into valid standard errors, with GPT-4o, DeBERTa-NLI, and PoliBERTweet as robustness labelers. We report the present results as our working estimates and will update them once validation is complete.

Selection of issues. The events are not a random sample. In the full design the estimand is defined over *Trump's communicated agenda*—his posting choice is the scope condition, not a bias—and the controversial/consolidated label is set ex ante from CES district preferences rather than from the observed response. The pilot uses the observed partisan stance distribution as a transparent proxy.

Identification. We will flesh out the quasi-random-timing argument with event-study pre-trend tests, placebo cutoffs, and a media common-shock control (and, where available, journalist/media-account timing showing Trump leads). The narrow ±7-day window guards against coincident shocks; cross-event spillover and concurrent events are handled with a stacked event-study.

Extensions. A panel VAR over persistent topics and frames (who leads whom over the whole period), a rhetoric-versus-votes consistency analysis linking tweet stance to matched roll-call votes (Gentzkow et al., 2019), and within-GOP heterogeneity (MAGA vs. establishment) are the natural next steps.

8 Conclusion

When a dominant national figure posts, Congress responds—but the response is asymmetric and issue-specific, and it splits cleanly along the two traps the title names. On a consolidated party agenda both parties mobilize and the in-party aligns. On a controversial signal, *strategically loyal Republicans* go quiet—only the least cross-pressured speak, and they align—while *Democrats face a prisoner's dilemma*, engaging hardest exactly where loud opposition most amplifies Trump's preferred framing. Reading congressional silence as a strategic choice, rather than as missing data, is essential to seeing both. The pilot establishes the measurement pipeline and the substantive pattern; the confirmatory study will rest on validated stance measurement, ex ante issue selection, the companion panel VAR, and the full identification battery.

References

Almond, D. and Du, X. (2020). Later bedtimes predict president trump's performance. *Economics Letters*.

Ansolabehere, S. and Kuriwaki, S. (2022). Congressional representation: Accountability from the constituent's perspective. *American Journal of Political Science*, 66(1):123–139.

Ansolabehere, S. and Kuriwaki, S. (2025). Collective representation in congress. *Perspectives on Politics*.

Arnold, R. D. (1990). *The Logic of Congressional Action*. Yale University Press.

Barberá, P., Casas, A., Nagler, J., Egan, P. J., Bonneau, R., Jost, J. T., and Tucker, J. A. (2019). Who leads? who follows? measuring issue attention and agenda setting by legislators and the mass public using social media data. *American Political Science Review*, 113(4):883–901.

Benoit, K., DeMarchi, S., Laver, C., Laver, M., and Ma, J. (2025). Using large language models to analyze political texts through natural language understanding. Working paper.

Burnham, M. (2025). Stance detection: A practical guide to classifying political beliefs in text. *Political Science Research and Methods*, 13(3):611–628.

Canes-Wrone, B., Brady, D. W., and Cogan, J. F. (2002). Out of step, out of office: Constituency and the electoral accountability of u.s. house members. *American Political Science Review*, 96(1):127–140.

Canes-Wrone, B., Herron, M. C., and Shotts, K. W. (2001). Leadership and pandering: A theory of executive policymaking. *American Journal of Political Science*, 45(3):532–550.

Carson, J. L., Koger, G., Lebo, M. J., and Young, E. (2010). The electoral costs of party loyalty in congress. *American Journal of Political Science*, 54(3):598–616.

Clinton, J. D. (2006). Representation in congress: Constituents and roll calls in the 106th house. *The Journal of Politics*, 68(2):397–409.

Cohen, J. E. (1995). Presidential rhetoric and the public agenda. *American Journal of Political Science*, 39(1):87–107.

Egami, N., Hinck, M., Stewart, B. M., and Wei, H. (2024). Using large language model annotations for the social sciences: A general framework of using predicted variables in downstream analyses. Working paper; design-based supervised learning (DSL); R package `dsl` (dsl.software).

Eshbaugh-Soha, M. and Peake, J. S. (2011). *Breaking Through the Noise: Presidential Leadership, Public Opinion, and the News Media*. Studies in the Modern Presidency. Stanford University Press.

Faris, R., Roberts, H., Etling, B., Bourassa, N., Zuckerman, E., and Benkler, Y. (2017). Partisanship, propaganda, and disinformation: Online media and the 2016 u.s. presidential election. Technical report, Berkman Klein Center for Internet & Society, Harvard University.

Fenno, R. F. (1973). *Congressmen in Committees*. Little, Brown.

Fouirnaies, A. and Hall, A. B. (2022). How do electoral incentives affect legislator behavior? evidence from u.s. state legislatures. *American Political Science Review*.

Gentzkow, M., Shapiro, J. M., and Taddy, M. (2019). Measuring group differences in high-dimensional choices: Method and application to congressional speech. *Econometrica*, 87(4):1307–1340.

Graham, M. H. and Svolik, M. W. (2020). Democracy in america? partisanship, polarization, and the robustness of support for democracy in the united states. *American Political Science Review*, 114(2):392–409.

Hausman, C. and Rapson, D. S. (2018). Regression discontinuity in time: Considerations for empirical applications. *Annual Review of Resource Economics*, 10:533–552.

Hemphill, L., Otterbacher, J., and Shapiro, M. A. (2013). What's congress doing on twitter? In *Proceedings of the 2013 Conference on Computer Supported Cooperative Work (CSCW)*.

Hirano, S., Snyder, Jr., J. M., Ansolabehere, S., and Hansen, J. M. (2010). Primary elections and partisan polarization in the u.s. congress. *Quarterly Journal of Political Science*, 5(2):169–191.

Kawintiranon, K. and Singh, L. (2022). PoliBERTweet: A pre-trained language model for analyzing political content on twitter. In *Proceedings of the Language Resources and Evaluation Conference (LREC)*.

King, G., Lam, P., and Roberts, M. E. (2017). Computer-assisted keyword and document set discovery from unstructured text. *American Journal of Political Science*, 61(4):971–988.

Laurer, M., van Atteveldt, W., Casas, A., and Welbers, K. (2024). Less annotating, more classifying: Addressing the data scarcity issue of supervised machine learning with deep transfer learning and BERT-NLI. *Political Analysis*, 32(1):84–100.

Lebo, M. J. and O'Geen, A. J. (2011). The president's role in the partisan congressional arena. *Journal of Politics*, 73(3):718–734.

Levitsky, S. and Ziblatt, D. (2018). *How Democracies Die*. Crown.

Mayhew, D. R. (1974). *Congress: The Electoral Connection*. Yale University Press.

Miller, W. E. and Stokes, D. E. (1963). Constituency influence in congress. *American Political Science Review*, 57(1):45–56.

Noble, B. S. (2023). Presidential cues and the nationalization of congressional rhetoric, 1973–2016. Working paper / manuscript.

Shepsle, K. A. (1972). The strategy of ambiguity: Uncertainty and electoral competition. *American Political Science Review*, 66(2):555–568.

Vishwanath, A. (2026). Value representation in the US congress. Working paper, Boston University.