

House of the Rising Sun

Strategically Loyal Republicans and Prisoner's-Dilemma Democrats

Da Gong (SUNY Geneseo) Andong Yan (HKU) Cantay Caliskan (U Rochester) Ezgi Kibris
(RIT) Eunseong Oh (UC Riverside)

July 8 EIP 2026

Outline

- 1 Motivation & Framework
- 2 Data
- 3 Empirical Strategy
- 4 Data Pipeline
- 5 Stance Results
- 6 Next Steps

Motivation: A Cult of Personality and the Separation of Powers

- **Democratic-backsliding concern:** a cult of personality around a party leader erodes the separation of powers — the legislature stops *checking* the executive and instead *defers* to the leader (Levitsky & Ziblatt 2018; Graham & Svolik 2020).
- **One channel is representation.** A member can fail to reflect local constituents for two distinct reasons:
 - **Loyalty/career channel:** the member fears punishment from party leadership for deviating from the party line.
 - **Voter-preference channel:** constituents themselves come to treat the leader's signal as dominant over their own issue preferences.
- **Setting:** Trump is an unprecedentedly salient figure; his Truth Social posts issue high-salience signals that rapidly structure elite attention and partisan messaging.
- **Empirical opportunity:** a stream of *datable, leader-issued signals*.

Central question

When a salient party leader shifts position, how do members of Congress respond — and to what extent does the response depend on the *leader's cue* versus *constituent preferences*?

What Counts as Representation Failure — and Our Assumptions

Normative benchmark: the trustee

- A member should exercise independent judgment in service of constituents' *interests*, not merely transmit their *will* — least of all a low-information affective cue (Burkean interest-vs-will: Miller & Stokes 1963; Canes-Wrone et al. 2001; Arnold 1990).
- Abandoning a district's issue preference to follow the presidential cue is a representation *failure* — and a failure *whichever* channel drives it (voter demand *or* loyalty pressure). Both substitute the leader's signal for the representative's job.

Deliberately weak assumptions

- Voters are rational *at the policy level* — but their preferences need not be good for them.
- We do *not* assume reps observe true constituent preferences (often infeasible; Vishwanath 2026).
- We assume only that observed behavior is a reelection-maximizing **equilibrium**.
- If voters are *not* even policy-rational, that only *deepens* the failure.

Design implication

Presidential-cue-following is a representation failure *regardless of channel* — which is what the empirical strategy can establish. Disentangling the channels is the next step.

Theory Foundation: The Loyalty–Accountability Trade-off

Strand 1: Strategic cue-taking

- Legislators use presidential cues strategically (Noble 2023).
- They *follow* supporter issue attention more than they lead it (Barberá et al. 2019).
- Effort flows where it is electorally productive (Fouirnaies & Hall 2022); reelection organizes the rest (Fenno 1973; Mayhew 1974).
- In-party rallies around a co-partisan president; the out-party resists (Lebo & O'Geen 2011; Cohen 1995; Eshbaugh-Soha & Peake 2011).

Strand 2: Electoral accountability

- Misalignment with the district reduces vote share (Canes-Wrone, Brady & Cogan 2002).
- Partisan loyalty on divisive issues is electorally costly (Carson et al. 2010; Hirano et al. 2010).

And a representation lens

- Classic position correspondence (Miller & Stokes 1963; Clinton 2006) → deeper values at scale (Vishwanath 2026; Ansolabehere & Kuriwaki 2022, 2025).

Resolution

Following the leader trades partisan coordination against the electoral cost of getting out of step with the district. **How a member resolves it depends on whether the leader's position is locally controversial.**

Controversial vs. Consolidated

- **Controversial:** opinion is internally divided — within the party's own electorate *and* between the party line and the district.
- **Consolidated:** the party and the district agree.
- **Proxy (pilot):** the within-party distribution of expressed stance.
 - **Tariffs:** Republicans near-unanimous (~84% support) ⇒ *consolidated* pole.
 - **Ukraine:** Republicans split roughly evenly (~49%) ⇒ *controversial* pole.

Two Asymmetric Traps

Republicans — strategic loyalty

- The leader is a co-partisan whose displeasure is electorally dangerous.
- On a *controversial* signal, a cross-pressured member cannot safely:
 - *align* $\Rightarrow$ alienate the district, or
 - *oppose* $\Rightarrow$ invite a primary.
- **Optimal move: go quiet.** Loyalty enforced through *silence*.
- Signature: in-party **under**-engagement.

Democrats — a prisoner's dilemma

- The signal comes from the out-party; the trap is *collective*.
- Each Democrat has a *private* incentive to attack a controversial Trump post (rallies the base, contests framing)...
- ... but loud, coordinated opposition **amplifies** the issue on Trump's preferred terms — collectively costly.
- Individually rational defection (engage) $\Rightarrow$ collectively worse outcome (amplify).
- Signature: out-party **over**-engagement.

Organizing theme

The two traps generate *opposite* engagement signatures on controversial issues: Republican silence vs. Democratic amplification.

Predictions

H1 — Controversial, extensive margin

On a controversial signal, **out-party (Democratic)** engagement jumps at the post while **in-party (Republican)** engagement does not — the signature of strategic loyalty/silence.

H2 — Consolidated, extensive margin

On a consolidated party-agenda signal, *both* parties' engagement jumps.

H3 — Intensive margin under selection

Conditional on engaging, Republicans who speak move toward Trump's position — concentrated among members whose districts impose the least local cross-pressure.

Two Behavioral Margins

- **Engagement / attention** (D_{ijt} , A_{ijt}): whether and how much a member speaks on the topic — the *extensive* margin.
- **Stance** (s_{ijt}): the position taken *conditional* on speaking — the *intensive* margin.

Research Questions

RQ1: Strategic response

When Trump issues a high-salience national signal, how do members adjust issue *attention* and rhetorical *stance*?

RQ2: National loyalty vs. local pressure

Do members follow the leader's cue even when it conflicts with local constituency opinion — and through which channel?

RQ3: Issue-specificity

Does the loyalty–accountability trade-off suppress in-party engagement *only* when the leader's position is locally controversial?

What we do

- Assemble a near-universe of original tweets by sitting MCs (2022–Jan 2026).
- Identify the precise Truth Social post that constitutes each Trump signal (T_e).
- Classify tweets for *relevance* and *stance* with open-weight LLMs (Qwen-2.5).
- Describe the partisan stance landscape across 24 signals.
- Estimate causal effects via **RD in event-time** around three sharp signals.

Pilot study; classifier validation and the full identification battery are in progress.

Contribution

- A scalable, reproducible LLM pipeline for elite attention and stance across $>2\text{M}$ candidate tweets.
- Timing-based identification of how a dominant national signaler moves congressional behavior.
- Evidence that strategic silence is **issue-specific** — reframing “silence” from missing data into a measurable strategic act.

The Sharp Signals (1/2): Ukraine — Two Controversial Shifts

Event 1 — Blame

ukraine_blame_01 Feb 19 2025, 15:47 UTC

“Think of it, a modestly successful comedian, Volodymyr Zelenskyy, talked the United States of America into spending \$350 Billion Dollars, to go into a War that couldn't be won, that never had to start, but a War that he, without the U.S. and 'TRUMP,' will never be able to settle. The United States has been SCAMMED, just like it has been SCAMMED by so many other Countries and 'Partners' for years. Zelenskyy, with no cards to play, is not the innocent man he portrays himself to be...”

Trump's position: **blame Ukraine / Zelensky** for prolonging the war; anti-Ukraine framing.

Event 2 — Reversal

ukraine_reversal_01 Sep 23 2025, 18:55 UTC

“After getting to know and fully understand the Ukraine/Russia Military and Economic situation and, after seeing the Economic trouble it is causing Russia, I think Ukraine, with the support of the European Union, is in a position to fight and WIN all of Ukraine back in its original form. With time, patience, and U.S. support, it can ALL be gotten back!”

Trump's position: **pro-Ukraine reversal** — endorses Ukraine's ability to win with EU + U.S. support.

*Both posts constitute sharp, unanticipated changes in Trump's publicly stated position on Ukraine — the **controversial** pole (~49% Republican support).*

The Sharp Signals (2/2): Tariffs — A Consolidated Signal

Event 3 — CCM tariffs

tariff_ccm_01 Feb 1 2025, 22:42 UTC

“Today, I have implemented a 25% Tariff on Imports from Mexico and Canada (10% on Canadian Energy), and a 10% additional Tariff on China. This was done through the International Emergency Economic Powers Act (IEEPA) because of the major threat of illegal aliens and deadly drugs killing our Citizens...”

Trump's position: **impose tariffs on Canada, Mexico, and China** — America First framing, justified on illegal immigration and fentanyl.

some statistics

The two Ukraine posts are *controversial* (Republicans split, ~49% support); the tariff announcement is a *consolidated* party-agenda signal (~84% Republican support).

Data Sources (1/2): Social Media & Signal Corpus

Congressional social-media panel

- All original tweets from sitting MCs, 2022–Jan 2026
- Full text + engagement (likes, RTs, replies, quotes, views)
- Linked to member via Bioguide crosswalk

Tweet filters applied:

- No retweets, replies, or quote-tweets
- English only; >20 characters

Trump signal corpus

- Truth Social (Feb 2022–Jan 2026)
- Original posts only; HTML stripped
- T_e identified per event via keyword search + researcher review

Stance classification

- Relevance: Qwen2.5-7B-Instruct (4-bit)
- Stance: Qwen2.5-14B-Instruct (4-bit)
- Robustness: GPT-4o validator + DeBERTaV3-NLI

Data Sources (2/2): District Covariates, CES & Roll-Call Votes

District-level covariates

- 2020 Trump / Biden vote share (district preference proxy)
- CES policy-preference estimates at district level (Ansolabehere & Kuriwaki 2022)
- Demographics: income, density, education, race

Member identity

- Bioguide crosswalk: party, state, chamber, district

Roll-call votes (rhetoric vs. votes)

- Congressional Record parser (Judd et al. 2017)
- Matched to CES decision agenda (Ansolabehere & Kuriwaki 2022)
- Links each vote to a policy topic j for text–vote comparison

Topic definition

- Topics defined *ex ante* from CES + Westwood, Grimmer & Hall (2025)
- Query expansion: King, Lam & Roberts (2017) keyword procedure

Two Identification Strategies & Behavioral Margins

Two sources of variation $\Rightarrow$ two complementary designs

- **Strategy 1 — persistent topics** (issues that simmer continuously) $\Rightarrow$ panel VAR: continuous co-movement of attention, stance, and framing. (*Estimates forthcoming.*)
- **Strategy 2 — novel posts as shocks** (sharp, datable switches) $\Rightarrow$ RD in event-time: discrete causal shocks. **This paper's estimates.**

Two behavioral margins

- **Extensive:** does the legislator engage with the topic at all? (D_{ijt})
- **Intensive:** conditional on engagement, what stance? (s_{ijt})

Empirical Setup: Notation & Definitions

Indices

- i — individual legislator (unique time series)
- $d(i)$ — congressional district of legislator i
- j — policy topic/issue
- t — day

Signals & preferences

- $s_{jt}^T \in [-1, +1]$ — Trump's daily stance on topic j (LLM-scored from Truth Social)
- **TrumpShare** $_{d(i)}$ — 2020 Trump vote share in district $d(i)$; proxy for district favor toward the national Trump agenda
- $\theta_{d(i),j}$ — CES-based district preference on issue j

Alignment gap:

$$\text{Gap}_{ijt} \equiv |s_{jt}^T - \theta_{d(i),j}|$$

Legislator outcomes

- $D_{ijt} \in \{0, 1\}$ — engagement indicator (*extensive margin*)
- $s_{ijt} \in [-1, +1]$ — expressed stance (*intensive margin*)
- A_{ijt} — daily attention share on topic j

Event-time (RD)

For Trump event e with timestamp T_e :

- Running variable $\tilde{d}_t = t - T_e$
- $1[t \geq T_e]$ — Trump raises or repositions on topic j
- Primary window: ± 7 days

FE structure: legislator α_i

Key assumption: $Y_{it}(0)$ continuous through T_e absent the post.

Strategy 1: Panel VAR Setup

Units (actors)

- **Trump** — single sender; Truth Social posts
- **Each legislator i** — individual MC time series
- **~20 media figures** — journalists / broadcasters (a common media shock could drive both Trump and Congress)

Two outcome series per actor $\times$ topic $j \times$ day t :

- A_{ijt} — attention share; log-odds $Z_{ijt}^A = \log\left(\frac{A_{ijt}}{1-A_{ijt}}\right)$
- $s_{ijt} \in [-1, +1]$ — expressed stance (Trump's series s_{jt}^T scored from Truth Social)

Bivariate VAR, $p = 7$ lags, topic FE α_j :

$$\begin{pmatrix} Z_{ijt}^A \\ s_{ijt} \end{pmatrix} = \alpha_j + \sum_{k=1}^7 \mathbf{B}_{ik} \begin{pmatrix} Z_{ijt-k}^A \\ s_{ijt-k} \end{pmatrix} + \varepsilon_{ijt}$$

Key IRF questions:

- Trump attention $\rightarrow$ legislator i attention: how fast, how large?
- Trump stance $\rightarrow$ legislator i stance: does s_{ijt} shift toward s_{jt}^T ?
- Heterogeneity: split by $\text{TrumpShare}_{d(i)}$, $\theta_{d(i),j}$, and party

Tracking stance as well as attention is what lets Strategy 1 speak to alignment, not merely agenda-setting.

Strategy 2: Regression Discontinuity in Time

Running variable: $\tilde{d}_t = t - T_e$ (event-relative day) **Cutoff:** T_e (Trump's Truth Social post)

Estimating equation


$$Y_{it} = \underbrace{\alpha_i}_{\text{member FE}} + \underbrace{\beta \cdot \mathbf{1}[t \geq T_e]}_{\text{within-member RD jump at } T_e} + \underbrace{\gamma \tilde{d}_t + \delta (\mathbf{1}[t \geq T_e] \cdot \tilde{d}_t)}_{\text{local-linear trend (each side)}} + \varepsilon_{it}$$

- **Two outcomes** estimated separately:
 - Y_{it}^{attn} : topic-tweet share of all tweets (defined $\forall$ member-days)
 - Y_{it}^{stance} : mean Qwen stance score (*conditional* on ≥ 1 topical tweet)
- β : within-member discontinuous jump at T_e — causal effect of the Trump signal
- Member FE α_i absorbed; window ± 7 days; SE clustered by `member_id`

Identification

Potential outcomes $Y_{it}(0)$ are *continuous* through T_e absent the post. Timing is plausibly quasi-random — Trump posts impulsively, off-schedule, often late at night (Almond & Du 2020). The estimand is robust to the post being endogenous in *origin*: a slow-moving upstream trigger is absorbed by the local trend.

Data Pipeline Overview



Stages 00b–01: Signal Identification & Event Registration

Stage 00b — Signal identification

- 1 Researcher defines seed keywords per event (e.g. “ukraine”, “zelensky”, “peace deal”)
- 2 Script searches Truth Social archive; ranks candidates by keyword match
- 3 Researcher reviews candidates; commits one post as T_e
- 4 Audit JSON written: *_te_selection.json

Stage 01 — Event registration

- T_e timestamp, window parameters, event type → events/events.parquet
- event_description and trump_position populated

Example: ukraine_blame_01

- T_e : 2025-02-19 15:47 UTC Type: stance_shift Pre/Post: 90/30 days
- Seed keywords: ukraine, zelensky, kyiv, peace, war, russia, dictator
- Selected post: *“Think of it, a modestly successful comedian, Volodymyr Zelenskyy, talked the United States of America into spending \$350 Billion Dollars, to go into a War that couldn’t be won, that never had to start...”*

Stages 02–03: Window Query & Dictionary Building

Stage 02 — Window query

- DuckDB query over Twitter Parquet archive
- Retrieves all original English tweets from member accounts in $[T_e - \text{pre}, T_e + \text{post}]$
- Typical yield: 90,000–120,000 tweets per event
- Output: candidates/{eid}_window.parquet

Stage 03 — LLM dictionary (GPU)

- Qwen2.5-7B-Instruct generates keyword dictionary from seed terms + event description
- Outputs: keywords, phrases, entities, hashtags, anchor context map
- Output: events/dictionaries/{eid}.json

Example: ukraine_blame_01 dictionary (excerpt)

Field	Terms
Keywords	ukraine, zelensky, kyiv, russia, peace, war, dictator, blame
Phrases	“peace deal”, “end the war”, “zelensky dictator”, “proxy war”
Hashtags	#Ukraine, #Zelensky, #StandWithUkraine
Entities	Kyiv, Zelenskyy, Putin, NATO, EU
Anchor	ukraine: [war, peace, russia, blame, support, negotiate]

Stages 04–05: Candidate Retrieval & Relevance Classification

Stage 04 — Candidate retrieval

- Window tweets matched against dictionary using anchor-aware contextual matching
- Typical yield: 2,000–25,000 candidates per event
- Output: `candidates/{eid}_candidates.parquet`

Stage 05 — Relevance (GPU)

- Qwen2.5-7B-Instruct, 4-bit, with checkpoint/resume
- Labels: 0 = not relevant, 1 = indirectly relevant, 2 = clearly about event
- Confidence via constrained single-token generation; checkpoint every 100 rows
- Output: `classified/{eid}_classified.parquet`

Relevance distribution: `ukraine_blame_01`

Label	Count	%
0 — not relevant	6,546	84.7%
1 — indirectly relevant	864	11.2%
2 — clearly relevant	326	4.2%
Total classified	7,736	
Relevant (1+2)	1,190	15.4%

Classified Tweets: Label 0 — Not Relevant (ukraine_blame_01, $n = 6,546$)

Interpretation

These tweets matched the Ukraine keyword dictionary but are *not* about Trump's blame-shift event — general foreign policy, domestic issues, MAGA messaging, etc. Retained in panel for attention-share denominator; excluded from stance analysis.

Handle	Conf	Tweet (excerpt)
@RepMcCaul	1.00	Biden's weakness invited Russian aggression. We need strong leadership to deter China next.
@GregoryMeeks	1.00	We must continue to support our allies. America's credibility is on the line worldwide.
@RepTonyGonzales	0.98	The border crisis is a national security threat. We have to stop the flow right now.

Classified Tweets: Label 1 — Indirectly Relevant (`ukraine_blame_01`, $n = 864$)

Interpretation

These tweets address U.S.–Ukraine policy, NATO support, or the Russia–Ukraine war in ways *tangentially* connected to Trump’s blame-shift — without explicitly referencing his Truth Social post.

Handle	Conf	Tweet (excerpt)
@RepMikeQuigley	0.84	We cannot abandon our allies. Cutting aid to Ukraine is a gift to Putin.
@BrianMast	0.79	The American people are tired of writing blank checks to foreign wars with no strategy.
@RepGregoryMeeks	0.71	The United States must remain a reliable partner to democracies fighting for their survival.

Classified Tweets: Label 2 — Clearly Relevant (`ukraine_blame_01`, $n = 326$)

Interpretation

These tweets explicitly respond to Trump blaming Ukraine/Zelensky for the war. They form the core sample for stance classification (Stage 06).

Handle	Conf	Tweet (excerpt)
@RepAdamSmith	0.97	Trump calling Zelensky a dictator is an insult to every democracy. Putin is the aggressor here.
@RepMariaSalazar	0.91	President Zelensky is fighting for freedom. America must stand with Ukraine, not blame the victim.
@RepScottPerry	0.88	Zelensky has refused every reasonable peace offer. The U.S. shouldn't keep funding a stalemate.

Stage 06: Stance Classification

Model & setup

- Qwen2.5-14B-Instruct (larger model for nuanced stance), 4-bit
- Input: all classified tweets (labels 0, 1, 2) from Stage 05 — every tweet has a stance score
- Constrained generation over {A, B, C, D} tokens

Label mapping	Token	Meaning	Numeric
	A	Oppose Trump's position (pro-Ukraine)	-1
	B	Neutral / no clear stance	0
	C	Support Trump's position (anti-Ukraine)	+1
	D	Unclear	99

Example: `ukraine_blame_01` Output: `classified/{eid}_classified.parquet`

	Dem	Rep
Oppose — pro-Ukraine (-1)	majority	minority
Neutral (0)		
Support — anti-Ukraine (+1)	minority	majority

Mean stance (all labels): -0.60 (std 0.61); strong partisan divergence in RD estimates.

Stages 07–09: Panel, Attention Metrics & Event Summary

Stage 07 — Daily panel

- Unit: member $i \times$ event $e \times$ day t
- `ukraine_share`: keyword-tweet share of all tweets that day
- `mean_stance_all`: daily mean stance across all classified tweets
- Output: `panel/panel.parquet`

Stages 08–09 — Attention metrics & summary

- Pre/post daily means, attention ratio, aggregate counts
- Output: `panel/attention_metrics.parquet`, `panel/event_summary.parquet`

Pipeline completion: 22 events

Stage	n Done	GPU
02 Window	22/22	No
03 Dictionary	21/22	Yes
04 Candidates	22/22	No
05 Relevance	22/22	Yes
06 Stance	22/22	Yes
07–09 Panel	22/22	No

The Partisan Stance Landscape

Key descriptive findings

- Democrats oppose Trump's position in **58–97%** of relevant tweets across all events
- Republicans support Trump in **23–84%** of tweets — wide within-party variation
- Partisan gap is large and consistent — parties rarely converge
- **Ukraine** (Rep: 49% support) and **NATO threat** (Rep: 43%) are the clearest exceptions to GOP alignment
- **Tariffs** shows the highest GOP support at 83.7%

Average stance across all 22 events

Metric	Democrats	Republicans
Avg % oppose	>80%	~18%
Avg % neutral	~12%	~15%
Avg % support	<5%	~65%
Highest support	7.3%	83.7%
Lowest support	0.0%	22.6%


Sample: ~60,000 classified tweets (conf ≥ 0.70), 22 events.

The consolidated→controversial spectrum is exactly the cross-issue variation the theory says governs whether in-party members can speak freely.

Stance: trump_tariffs_01 (Mar 3 2025 — Liberation Day)

Partisan distribution

	Dem	Rep
% Oppose	91.3	4.8
% Neutral	5.9	8.6
% Support	2.3	83.7



takeaway

Tariffs: highest GOP alignment in the sample (83.7% support).


Contrast with Ukraine (49% R support): tariffs are a *consolidated* Republican position, where the loyalty-accountability trade-off imposes minimal electoral cost.

Democrats near-unanimously opposed (91.3%), consistent with the counter-signaling incentive.

Stance: ukraine_blame_01 (Feb 19 2025)

Partisan distribution

	Dem	Rep
% Oppose	96.3	41.7
% Neutral	2.3	9.0
% Support	0.6	48.9





takeaway

Republican support (49%) is the lowest in the sample — nearly split. Ukraine is *not* a fully consolidated GOP position.

Democrats are near-unanimously opposed (96%) — the lowest Democrat support of any event.

Stance Validation: Ukraine Events



Left: ukraine_blame_01 (Feb 19 2025) **Right:** ukraine_reversal_01 (Sep 23 2025)

Distribution of Qwen stance scores (-1 = oppose/pro-Ukraine, 0 = neutral, $+1$ = support/anti-Ukraine) by party.

RD Results: Discontinuity at T_e (Member FE)

Spec: $Y_{it} = \alpha_i + \beta \mathbf{1}[t \geq T_e] + \gamma \tilde{d}_t + \delta (\mathbf{1}[t \geq T_e] \cdot \tilde{d}_t) + \varepsilon_{it}$


Member FE absorbed; SE clustered by member. ** $p < 0.05$, *** $p < 0.01$

Event / Outcome	$\hat{\beta}$	SE	p	N
<i>ukraine_blame_01</i> — <i>Trump blames Ukraine/Zelensky (Feb 19 2025)</i>				
Attention share	+0.030	0.014	0.037**	3,922
Stance (Qwen)	-0.075	0.083	0.370	803
<i>ukraine_reversal_01</i> — <i>Trump shifts pro-Ukraine (Sep 23 2025)</i>				
Attention share	+0.030	0.011	0.007***	3,911
Stance (Qwen)	+0.252	0.153	0.102	589

Takeaway


Both events produce a significant within-member jump in Ukraine *attention share* at T_e . Pooled stance RD is in the expected direction but underpowered — improving LLM accuracy is the priority.

RD Visual: ukraine_blame_01 (Feb 19 2025)



Scatter = within-member demeaned daily means. Lines = local-linear fit on demeaned outcomes (pre: ● gold; post: ● green). Within-member jump at T_e : $\hat{\beta} = +0.030$, $se = 0.014$, $p = 0.037$.

RD Visual: ukraine_reversal_01 (Sep 23 2025)



Scatter = within-member demeaned daily means. Lines = local-linear fit on demeaned outcomes (pre: ● gold; post: ● green). Within-member jump at T_e : $\hat{\beta} = +0.030$, $se = 0.011$, $p = 0.007$.

RD Heterogeneous Effects: Attention by Party & Local Issue

Outcome: Ukraine attention share. Member FE absorbed; SE clustered by member.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Subgroup	RD jump	SE	p	N
<i>ukraine_blame_01 (Feb 19 2025)</i>				
Republican (all)	+0.020	0.019	0.315	1,893
Democrat (all)	+0.042	0.021	0.051*	2,029
D Low Trump	+0.040	0.021	0.059*	1,985
<i>ukraine_reversal_01 (Sep 23 2025)</i>				
Republican (all)	+0.011	0.015	0.463	1,873
Democrat (all)	+0.047	0.016	0.004***	2,038
D Low Trump	+0.054	0.016	0.001***	1,997
D (Low T, High Local)	+0.060	0.019	0.002***	1,286

Pattern (H1 + H4)

Democrats drive both RD jumps; Republicans show no response (strategic silence). On reversal the Democratic effect is *strongest* exactly where amplification is most costly to the party — Low-Trump, High-Local (pro-Ukraine) districts (+6.0pp, $p = 0.002$): the prisoner's-dilemma signature.

RD Heterogeneous Effects: Stance by Party & Local Issue


Outcome: mean Qwen stance score (-1 = pro-Ukraine/oppose Trump; $+1$ = anti-Ukraine/support Trump).
Member FE absorbed; SE clustered by member. * $p < 0.10$, ** $p < 0.05$

Subgroup	RD jump	SE	p	N
<i>ukraine_blame_01 (Feb 19 2025)</i>				
R (all)	-0.208	0.097	0.035**	386
R High Trump	-0.186	0.099	0.065*	365
R (High Trump, Low Local)	-0.279	0.144	0.058*	235
D (all)	+0.104	0.130	0.425	417
<i>ukraine_reversal_01 (Sep 23 2025)</i>				
R (all)	+0.463	0.226	0.045**	244
R High Trump	+0.327	0.203	0.111	224
R (High Trump, Low Local)	+0.366	0.300	0.230	138
D (all)	+0.087	0.200	0.663	345

Takeaway (H3)

Engaged Republicans *do* align with Trump when they tweet. High-Trump, Low-Local R districts amplify his blame framing (-0.21 , $p = 0.035$) and follow his pro-Ukraine reversal ($+0.46$, $p = 0.045$). Democrats hold their pro-Ukraine stance throughout.

RD Visual: tariff_ccm_01 (Feb 1 2025)



Scatter = within-member demeaned daily means (tariff-topic tweet share). Lines = local-linear fit on demeaned outcomes (pre: ● gold; post: ● green). Vertical line = T_e (CCM tariff announcement, Feb 1 2025). Both parties show a clear upward jump at T_e .

RD Heterogeneous Effects: Attention — tariff_ccm_01

Outcome: tariff-topic tweet share. Member FE absorbed; SE clustered by member.

** $p < 0.05$, *** $p < 0.01$

Subgroup	RD jump	SE	p	N
Republican (all)	+0.130	0.030	0.000***	2,026
R High Trump	+0.132	0.032	0.000***	1,868
R High Trump + Pro-Tariff Local	+0.118	0.042	0.006***	1,115
R High Trump + Anti-Tariff Local	+0.131	0.052	0.014**	701
R Low Trump	+0.119	0.076	0.162	97
Democrat (all)	+0.074	0.020	0.000***	3,255
D Low Trump	+0.098	0.025	0.000***	2,094
D Low Trump + Anti-Tariff Local	+0.115	0.029	0.000***	1,261
D Low Trump + Pro-Tariff Local	+0.064	0.046	0.168	808

Pattern (H2)

Both parties significantly increase engagement. Republicans in High-Trump districts respond uniformly regardless of local tariff sentiment (+13pp either way). Anti-Tariff Democrats also spike (+11.5pp) — amplifying opposition.

RD Heterogeneous Effects: Stance — tariff_ccm_01

Outcome: mean Qwen stance score (-1 = anti-tariff/oppose Trump; $+1$ = pro-tariff/support Trump).
Member FE absorbed; SE clustered by member. *** $p < 0.01$

Subgroup	RD jump	SE	p	N
Republican (all)	+0.179	0.059	0.003***	714
R High Trump	+0.179	0.062	0.005***	665
R High Trump + Pro-Tariff Local	+0.227	0.069	0.002***	421
R High Trump + Anti-Tariff Local	+0.105	0.138	0.450	218
R Low Trump	+0.147	0.163	0.401	32
Democrat (all)	+0.067	0.078	0.387	800
D Low Trump	+0.049	0.121	0.685	432
D Low Trump + Anti-Tariff Local	+0.068	0.164	0.682	248

Takeaway

Republicans align with Trump's pro-tariff stance ($+0.18$, $p = 0.003$); strongest in High-Trump, Pro-Tariff districts ($+0.23$, $p = 0.002$). **Democrats show no significant stance shift** — they increase engagement to oppose, but their mean stance stays firmly anti-tariff.

Contrasting Issues (1/2): Ukraine — Controversial

Ukraine — Controversial signal

49% Republican support

GOP trade-off: R districts are split 49/41 for/against Ukraine $\Rightarrow$ engaging publicly imposes electoral risk.

Attention RD

- **Democrats:** significant jump (+4.2pp, $p = 0.051$) on Blame; (+4.7pp, $p = 0.004$) on Reversal — counter-signaling raises salience
- **Republicans:** no significant response in either event
- Low-Trump D districts lead; effect strongest where amplification is costliest

Stance RD

- Engaged **R** (High-Trump, Low-Local) *do* shift toward Trump's frame: Blame -0.21 ($p = 0.035$); Reversal $+0.46$ ($p = 0.045$)
- **D** hold pro-Ukraine stance throughout — no shift in either event

Takeaway

Cross-pressured Republicans go quiet; those who *do* tweet align with Trump. Democrats amplify hardest exactly where loud opposition is most costly — the prisoner's dilemma.

Contrasting Issues (2/2): Tariffs — Consolidated

Tariffs — Consolidated party agenda 84% Republican support (highest in sample)

No GOP trade-off: Tariffs = MAGA/America First core agenda $\Rightarrow$ small local-accountability penalty for engaging.

Attention RD

- **Republicans:** large, uniform jump (+13pp, $p < 0.001$) — party agenda-setting; holds regardless of local tariff sentiment (+11.8pp pro-local, +13.1pp anti-local)
- **Democrats:** also jump significantly (+7.4pp, $p < 0.001$) to amplify opposition

Stance RD

- **R** shift pro-tariff (+0.18, $p = 0.003$); strongest in High-Trump, Pro-Tariff districts (+0.23, $p = 0.002$)
- **D:** engage more but stance does not shift — remains firmly anti-tariff

Takeaway

Republicans mobilize fully on consolidated agenda items (engagement *and* stance follow Trump). Democrats consistently counter-signal.

Limitations & Next Steps

Stance measurement & validation

- Off-the-shelf accuracy $\approx 70\%$; error attenuates stance toward zero (Burnham 2025).
- Human gold standard + codebook; stronger production classifier; *design-based supervised learning* (Egami et al. 2024) for valid SEs; GPT-4o, DeBERTa-NLI, PoliBERTweet robustness labelers.

Selection of issues

- Estimand defined over *Trump's communicated agenda*; controversial/consolidated set *ex ante* from CES.

Identification

- Event-study pre-trend tests, placebo cutoffs, media common-shock control.
- ± 7 -day window guards against coincident shocks; stacked event-study for spillover.

Extensions

- Panel VAR over persistent topics (who leads whom).
- Rhetoric-vs-votes consistency (Gentzkow et al. 2019).
- Within-GOP heterogeneity (MAGA vs. establishment).
- Disentangle the channels: voter demand vs. loyalty pressure.

Thank you

Questions & Comments Welcome

Da Gong
dgong005@ucr.edu

Appendix: Full Data Funnel I

Pipeline data funnel: window → candidates → classified. *pre_days* and *post_days* give window half-widths around T_e .

event_id	type	Te	pre_days	post_days	window_start	wind
biden_dropout_01	novel_shock	2024-06-21	90	30	2024-03-23	2024
charlie_kirk_assassination_01	novel_shock	2025-08-27	90	30	2025-05-29	2025
dei_executive_order_01	stance_shift	2025-01-19	90	30	2024-10-21	2025
doge_disbanded_01	stance_shift	2025-11-20	90	30	2025-08-22	2025
doge_establish_01	novel_shock	2024-12-04	90	30	2024-09-05	2025
government_shutdown_01	stance_shift	2025-03-13	90	30	2024-12-13	2025
greenland_01	novel_shock	2024-12-22	30	30	2024-11-22	2025
greenland_02	novel_shock	2026-01-14	30	30	2025-12-15	2026
haitian_pets_01	novel_shock	2024-08-15	90	30	2024-05-17	2024
minneapolis_ice_shooting_01	novel_shock	2025-12-11	90	30	2025-09-12	2026
nato_threat_01	stance_shift	2025-03-08	90	30	2024-12-08	2025
roe_wade_01	stance_shift	2022-06-24	90	30	2022-03-26	2022
trump_assassination_attempt_01	novel_shock	2024-07-01	90	30	2024-04-02	2024
trump_convicted_01	novel_shock	2024-05-31	90	30	2024-03-02	2024
trump_musk_feud_01	stance_shift	2025-06-05	90	30	2025-03-07	2025
trump_tariffs_01	novel_shock	2025-03-03	90	30	2024-12-03	2025

Appendix: Full Attention Metrics I

Attention metrics: pre- and post- T_e daily tweet means and ratio.

Event	Pre mean	Post mean	Ratio	Confirmed	Pre flat
greenland_01	0.833000	2.581000	3.097000	1	0
greenland_02	4.300000	6.581000	1.530000	0	0
biden_dropout_01	18.789000	48.194000	2.565000	1	0
charlie_kirk_assassination_01	11.900000	65.839000	5.533000	1	0
dei_executive_order_01	2.711000	17.936000	6.616000	1	0
doge_disbanded_01	15.700000	13.387000	0.853000	0	0
doge_establish_01	6.522000	23.548000	3.610000	1	0
government_shutdown_01	47.122000	50.161000	1.064000	0	0
haitian_pets_01	2.267000	1.871000	0.825000	0	0
minneapolis_ice_shooting_01	54.267000	49.194000	0.906000	0	0
nato_threat_01	4.367000	3.936000	0.901000	0	0
roe_wade_01	1.911000	4.516000	2.363000	1	0
trump_assassination_attempt_01	5.844000	16.677000	2.854000	1	0
trump_convicted_01	7.578000	11.323000	1.494000	0	0
trump_musk_feud_01	24.467000	40.484000	1.655000	0	0
trump_tariffs_01	56.144000	110.871000	1.975000	0	0

Appendix: Event Summary (Stage 09) I

Event-level summary statistics for all events (Stage 09 output). *attention_confirmed* = True when $\text{post-}T_e \text{ mean} \geq 2 \times \text{pre-}T_e \text{ mean}$.

event_id	event_type	event_date	issue_domain	confounds_noted	attention_ratio	attention_c
greenland_01	novel_shock	2024-12-22	foreign_policy	—	3.259300	
greenland_02	novel_shock	2026-01-14	foreign_policy	—	0.741600	